

# Week 6: Tree Ensembles Assignment

## ARIMA/GARCH + Random Forest & XGBoost for Portfolio Optimization

MSc Banking and Finance – FinTech Course

**Due: we will discuss this in class..**

## 1 Assignment Overview

**Objective:** Implement tree-based ensemble methods for portfolio optimization using Week 5 engineered features.

**Deliverables:**

- Jupyter Notebook (.ipynb) with code and analysis
- PDF Report (3-5 pages) summarizing findings
- Results file (.csv) with portfolio weights

## 2 Assignment Tasks

### 2.1 Task 1: Model Implementation & Comparison (30%)

**Requirements:**

1. Implement three models: Random Forest, Gradient Boosting, and XGBoost
2. Train each model on Week 5 features (use at least 15 features)
3. Perform 5-fold cross-validation for each model
4. Report MSE,  $R^2$ , and MAE for train and test sets
5. Conduct paired t-tests between models (report p-values)

**Expected Output:**

Model	Train $R^2$	Test $R^2$	MSE	p-value
Random Forest	0.XXX	0.XXX	0.XXXXXXX	—
Gradient Boosting	0.XXX	0.XXX	0.XXXXXXX	0.XXX
XGBoost	0.XXX	0.XXX	0.XXXXXXX	0.XXX

### 2.2 Task 2: Hyperparameter Tuning (25%)

**Requirements:**

1. Select your best-performing model from Task 1
2. Define hyperparameter grid with at least 5 parameters
3. Use GridSearchCV with 5-fold cross-validation
4. Create validation curves for top 3 hyperparameters
5. Report optimal hyperparameters with justification

**Minimum Grid Size:**

- Random Forest: Test at least 36 combinations
- Gradient Boosting: Test at least 36 combinations
- XGBoost: Test at least 48 combinations

**Justification:** Explain why optimal values make sense (150-200 words).

## 2.3 Task 3: Feature Importance & SHAP (25%)

### Requirements:

1. Extract built-in feature importance from your best model
2. Compute SHAP values for test set
3. Create SHAP summary plot
4. Compare built-in importance vs. SHAP importance (plot both)
5. Identify top 5 features and explain their impact (200-250 words)

### Analysis Questions:

- Which features drive predictions most strongly?
- Do you observe any non-linear interactions?
- Are results financially interpretable?

## 2.4 Task 4: Portfolio Optimization & Backtesting (20%)

### Requirements:

1. Use best model to generate portfolio weights for BTC, ETH, DOGE
2. Implement rolling-window backtesting (minimum 30-day window)
3. Compare three strategies:
  - Equal-weighted (1/3, 1/3, 1/3)
  - Week 5 Lasso-based portfolio
  - Week 6 Tree-based portfolio
4. Calculate: Sharpe ratio, max drawdown, win rate, total return
5. Visualize cumulative returns

### Performance Table:

Strategy	Sharpe	Max DD	Win Rate	Total Ret.
Equal-Weighted	X.XX	-XX.X%	XX.X%	XX.X%
Lasso (Week 5)	X.XX	-XX.X%	XX.X%	XX.X%
Tree Ensemble	X.XX	-XX.X%	XX.X%	XX.X%

## 3 Bonus Tasks (Optional, +10%)

Choose **one** of the following:

### 3.1 Option A: Custom Ensemble Stacking

Combine Ridge (Week 5), Random Forest, and XGBoost predictions using optimal weights. Show improvement over individual models.

### 3.2 Option B: Regime Detection

Use decision tree splits to identify market regimes. Create separate portfolio strategies for each regime.

### 3.3 Option C: Early Stopping Analysis

Implement early stopping for Gradient Boosting/XGBoost. Plot training vs. validation error curves. Determine optimal number of iterations.

## 4 Evaluation Rubric

Criterion	Description	Points
<b>Code Quality</b>	Clean, documented, functional, follows best practices	20
<b>Model Performance</b>	All models trained, tuned, significant improvement shown	25
<b>Analysis Depth</b>	Insightful interpretations, connects theory to results	25
<b>Statistical Rigor</b>	Proper testing, cross-validation, significance tests	15
<b>Visualizations</b>	Professional, clear, informative plots	10
<b>Report Quality</b>	Well-written, organized, concise	5
<b>Total</b>		<b>100</b>
<b>Bonus</b>	(Optional)	<b>+10</b>

## 5 Quick Start Code Templates

### 5.1 Model Training

```

1 # Random Forest
2 from sklearn.ensemble import RandomForestRegressor
3 rf = RandomForestRegressor(n_estimators=100, max_depth=10,
4                             random_state=42, n_jobs=-1)
5 rf.fit(X_train, y_train)
6
7 # XGBoost
8 import xgboost as xgb
9 xgb_model = xgb.XGBRegressor(n_estimators=100, learning_rate=0.1,
10                               max_depth=6, reg_alpha=0.1, reg_lambda=1.0)
11 xgb_model.fit(X_train, y_train)

```

### 5.2 Hyperparameter Tuning

```

1 from sklearn.model_selection import GridSearchCV
2
3 param_grid = {
4     'n_estimators': [50, 100, 200],
5     'max_depth': [5, 10, 15],
6     'learning_rate': [0.01, 0.05, 0.1]
7 }
8
9 grid_search = GridSearchCV(model, param_grid, cv=5,
10                             scoring='neg_mean_squared_error')
11 grid_search.fit(X_train, y_train)
12 print(f"Best params: {grid_search.best_params_}")

```

### 5.3 SHAP Analysis

```

1 import shap
2
3 explainer = shap.TreeExplainer(model)
4 shap_values = explainer.shap_values(X_test)
5

```

```
6 # Summary plot
7 shap.summary_plot(shap_values, X_test, feature_names=feature_names)
```

## 6 Detailed Grading Breakdown

### 6.1 Excellent (90-100%)

- All models properly implemented with comprehensive analysis
- Rigorous hyperparameter tuning with clear justification
- Deep SHAP interpretation connecting to finance theory
- Professional visualizations with publication quality
- Statistical tests properly conducted and interpreted

### 6.2 Good (75-89%)

- Models trained with adequate tuning
- SHAP analysis present with reasonable interpretation
- Clear visualizations and proper documentation
- Some statistical testing conducted

### 6.3 Satisfactory (60-74%)

- Basic implementation functional
- Minimal tuning and analysis
- Adequate visualizations
- Limited statistical validation

## 7 Common Mistakes to Avoid

1. **Not using cross-validation** – Single train/test split is insufficient
2. **Scaling features for trees** – Trees don't need StandardScaler
3. **Default hyperparameters** – Always tune your models
4. **Ignoring overfitting** – Monitor train-test gap
5. **No SHAP analysis** – Feature importance alone is insufficient
6. **Missing statistical tests** – Must compare models rigorously
7. **Poor documentation** – Explain your code and decisions
8. **Late submission** – Start early, deadline is firm

## 8 Resources

### 8.1 Documentation

- scikit-learn: <https://scikit-learn.org/stable/>
- XGBoost: <https://xgboost.readthedocs.io/>
- SHAP: <https://shap.readthedocs.io/>