

# Week 6: Tree Ensembles for Portfolio Optimization

## Big Data & Statistical Learning for Finance

MSc Banking and Finance – FinTech Course

A.U.E.B.

October 14, 2025

# Outline

- 1 Introduction & Motivation
- 2 Decision Trees: Foundation
- 3 The Bias-Variance Tradeoff
- 4 Random Forest (Bagging)
- 5 Gradient Boosting
- 6 XGBoost: State-of-the-Art
- 7 Feature Importance & SHAP
- 8 Applying Trees to Portfolio Optimization
- 9 Hyperparameter Tuning
- 10 Model Comparison & Validation
- 11 Key Takeaways
- 12 Practical Advice
- 13 Detailed Code Examples
- 14 Connection to Course
- 15 Further Reading
- 16 Statistical Model Comparison

# Why Should You Care About Trees?

## The Problem with Linear Models

Linear models assume relationships are **straight lines**:

$$\hat{y} = w_0 + w_1x_1 + w_2x_2 + \dots + w_px_p$$

## Real Life is Complicated!

- Non-linear interactions
- Regime changes (bull vs. bear)
- Threshold effects

## Trees Excel At:

- Capturing non-linearities
- Feature interactions
- Robust to outliers
- No scaling needed!

# What Makes Finance Special?

## Financial Markets Are Complex

- **Regime Changes:** Markets behave differently when volatility is high vs. low
- **Interactions:** Good returns + low risk = great, but good returns + high risk = scary
- **Non-linear Relationships:** What works in bull markets fails in bear markets

**Trees naturally learn these complex rules!**

# How Decision Trees Work

## Like Playing 20 Questions!

A decision tree recursively splits data based on feature thresholds.

### Splitting Criterion (MSE):

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})^2$$

### Example: Predicting Bitcoin Returns

- Is volatility  $> 3\%$ ?
  - **Yes**: Risky!  $\rightarrow$  Predict  $-2\%$
  - **No**: Safe!  $\rightarrow$  Predict  $+3\%$

Choose split that minimizes weighted MSE:

$$\text{MSE}_{\text{total}} = \frac{|D_L|}{n} \text{MSE}_L + \frac{|D_R|}{n} \text{MSE}_R$$

# Single Tree: Pros and Cons

## Advantages

- Captures non-linearities
- No feature scaling needed
- Handles mixed data types
- Interpretable (visualizable)

## Disadvantages

- **High variance**
- Overfitting prone
- Unstable predictions
- Sensitive to data changes

**Solution: Use Ensembles!**

# Bias-Variance Decomposition

## Expected Prediction Error

$$\mathbb{E}[(y - \hat{f}(x))^2] = \underbrace{\text{Bias}[\hat{f}(x)]^2}_{\text{Underfitting}} + \underbrace{\text{Var}[\hat{f}(x)]}_{\text{Overfitting}} + \underbrace{\sigma_\epsilon^2}_{\text{Irreducible}}$$

### Bias (Systematic Error):

- Model too simple
- Missing patterns
- Always wrong in same way

### Variance (Random Error):

- Model too complex
- Overfits training data
- Unpredictable on new data

# The Goldilocks Principle

Model Type	Bias	Variance
Simple (Linear)	High	Low
Single Tree	Low	High
<b>Ensemble of Trees</b>	Low	Low

## Key Insight

**Ensembles** = Many trees working together = Best of both worlds!



# Wisdom of the Crowd

## Think of it Like This...

**Scenario:** Estimate how many jellybeans are in a jar.

- **Bad Strategy:** Ask one person (might be way off)
- **Good Strategy:** Ask 100 people and average their guesses

**Result:** The average is usually very close to the true number!

## Random Forest = Wisdom of the Crowd for Trees

- ① Train many trees on different random data samples
- ② At each split, only consider a random subset of features
- ③ Average all predictions

# Random Forest Algorithm

---

## Algorithm 1 Random Forest

---

- 1: **Input:** Dataset  $D$ , number of trees  $B$
  - 2: **Parameter:** Features per split  $m < p$
  - 3: **for**  $b = 1$  to  $B$  **do**
  - 4:   Draw bootstrap sample  $D_b$  from  $D$
  - 5:   Train tree  $T_b$  on  $D_b$  where:
  - 6:     At each split, randomly select  $m$  features
  - 7:     Choose best split among the  $m$  features
  - 8: **end for**
  - 9: **Output:**  $\hat{y}_{\text{RF}}(x) = \frac{1}{B} \sum_{b=1}^B T_b(x)$
-

# Why Random Forest Works Better

## Variance Reduction Through Averaging

For  $B$  independent predictions with variance  $\sigma^2$ :

$$\text{Var} \left( \frac{1}{B} \sum_{b=1}^B Z_b \right) = \frac{\sigma^2}{B}$$

As  $B \rightarrow \infty$ , variance  $\rightarrow 0$

## The Magic of Randomness

- **Bootstrap sampling:** Each tree sees different data
- **Feature sampling:** Trees make different mistakes
- **Averaging:** Mistakes cancel out!

**Mathematical fact:** Averaging 100 independent predictions reduces error by  $\approx 10$ -fold!

# Key Hyperparameters

Parameter	Description	Typical Values
n_estimators	Number of trees	50–500
max_depth	Tree depth	10–20 or None
min_samples_split	Min samples to split	2–20
max_features	Features per split	'sqrt', 'log2'
bootstrap	Use bootstrap sampling	True

## Free Validation Trick: Out-of-Bag (OOB) Score

About 1/3 of data isn't used by each tree  $\implies$  use for validation!

$$P(\text{obs } i \notin \text{bootstrap}) = \left(1 - \frac{1}{n}\right)^n \approx e^{-1} \approx 0.368$$

# Learning from Mistakes

## Random Forest

- Trees trained **in parallel**
- Independent of each other
- Average predictions

## Gradient Boosting

- Trees trained **sequentially**
- Each fixes previous mistakes
- Additive model

## Basketball Free Throw Analogy

- ➊ **Attempt 1:** Miss – ball hits front of rim
- ➋ **Lesson:** Shoot with more power
- ➌ **Attempt 2:** Miss – now it's too far (back of rim)
- ➍ **Lesson:** Reduce power slightly
- ➎ **Attempt 3:** Adjust... closer!

# Gradient Boosting Process

## Step-by-Step

- 1 **Start simple:** Begin with average prediction
- 2 **Find mistakes:** Calculate residuals (errors)

$$r_i = y_i - \hat{y}_i$$

- 3 **Train tree on mistakes:** Build small tree to predict errors
- 4 **Update model:** Add correction

$$F_m(x) = F_{m-1}(x) + \eta \cdot h_m(x)$$

where  $\eta$  is the learning rate

- 5 **Repeat:** Go back to step 2

## Gradient Descent in Function Space

Minimize loss  $L(y, F(x))$  by iteratively moving in the negative gradient direction:

$$F_m = F_{m-1} - \eta \nabla_F L(y, F_{m-1})$$

For squared loss:  $L(y, \hat{y}) = (y - \hat{y})^2$

$$-\frac{\partial L}{\partial \hat{y}} = y - \hat{y} = \text{residual}$$

## Key Insight

Training a tree on residuals  $\approx$  Following negative gradient!

# Important Settings

Parameter	Description	Typical Values
<code>n_estimators</code>	Boosting rounds	100–1000
<code>learning_rate</code>	Shrinkage $\eta$	0.01–0.1
<code>max_depth</code>	Tree depth	3–7 (shallow!)
<code>subsample</code>	Row sampling	0.5–1.0

## Tuning Tips

- Lower `learning_rate` + more `n_estimators`  $\implies$  Better (but slower)
- Typical: `learning_rate`=0.01 with `n_estimators`=1000
- Keep trees shallow – each tree only fixes small part of problem
- `subsample` < 1.0 adds randomness (like bagging)



# What Makes XGBoost Special?

## eXtreme Gradient Boosting

Enhanced gradient boosting with three superpowers:

- ① **Built-in regularization** – Prevents overfitting automatically
- ② **Optimized for speed** – Uses computer tricks to run fast
- ③ **Smart missing data handling** – Doesn't break with missing values

## Why XGBoost Dominates Competitions

XGBoost is the “Swiss Army knife” of machine learning:

- Random Forest's column sampling (randomness)
- Gradient Boosting's sequential learning
- Ridge/Lasso regularization
- Plus computational tricks (GPU support, parallel processing)

# XGBoost Objective Function

## Regularized Objective

$$\mathcal{L}(\phi) = \sum_{i=1}^n \ell(y_i, \hat{y}_i) + \sum_{k=1}^K \Omega(f_k)$$

where  $\Omega(f)$  is the regularization penalty:

$$\Omega(f) = \gamma T + \frac{1}{2} \lambda \sum_{j=1}^T w_j^2$$

- $T$  = number of leaves
- $w_j$  = leaf weights
- $\gamma$  = complexity penalty (like L0 regularization)
- $\lambda$  = L2 regularization (like Ridge!)

**Effect:** Forces trees to stay simple and learn only important patterns.

# XGBoost in Python

```
import xgboost as xgb

xgb_model = xgb.XGBRegressor(
    n_estimators=100,          # Boosting rounds
    learning_rate=0.1,         # Shrinkage
    max_depth=6,               # Tree depth
    min_child_weight=1,        # Min sum of weights in leaf
    subsample=0.8,             # Row sampling
    colsample_bytree=0.8,      # Column sampling
    reg_alpha=0.1,             # L1 regularization
    reg_lambda=1.0,            # L2 regularization
    gamma=0,                   # Min loss reduction
    random_state=42
)

xgb_model.fit(X_train, y_train)
predictions = xgb_model.predict(X_test)
```

# Understanding What Your Model Learned

## Why This Matters

Just because your model works doesn't mean you should trust it blindly!  
You need to understand:

- Which features actually matter?
- Is the model using real patterns or random noise?
- Can you explain decisions to your boss/client?

## Three Methods:

- ① **Impurity-based:** Count splits (fast but biased)
- ② **Permutation importance:** Shuffle feature and measure drop (reliable)
- ③ **SHAP values:** Fair attribution based on game theory (best!)

# Permutation Importance

## The Cookie Baking Analogy

Want to know if sugar is important?

- 1 Bake cookies normally → Taste: 90/100
- 2 Randomly replace sugar amounts → Taste: 20/100
- 3 **Conclusion:** Sugar is VERY important (70 point drop!)

## Algorithm

For each feature  $j$ :

- 1 Measure model performance on original data
- 2 Randomly shuffle feature  $j$  values
- 3 Measure new performance
- 4 Importance = Performance drop

Bigger drop = More important feature!

# SHAP: Shapley Additive exPlanations

## What is SHAP?

Based on **Shapley values** from game theory – fairly divides prediction among features.

### Decompose prediction:

$$f(x) = \phi_0 + \sum_{j=1}^p \phi_j$$

- $\phi_0 = \mathbb{E}[f(X)]$  – expected prediction (baseline)
- $\phi_j$  – SHAP value for feature  $j$  (contribution)

### Example: Predicting Bitcoin Return

- Base prediction: +1.5% (average)
- Volatility is high today: -0.8% (pushes down)
- Momentum is positive: +1.2% (pushes up)
- **Final prediction:**  $1.5 - 0.8 + 1.2 = +1.9\%$

# SHAP Properties

## Fairness Axioms

SHAP values satisfy:

- 1 **Efficiency:**  $\sum_{j=1}^P \phi_j = f(x) - \mathbb{E}[f(X)]$
- 2 **Symmetry:** Equal contributions  $\implies$  equal SHAP values
- 3 **Dummy:** Irrelevant feature  $\implies$  SHAP value = 0
- 4 **Additivity:** For ensemble  $f = f_1 + f_2$ :  $\phi_j^f = \phi_j^{f_1} + \phi_j^{f_2}$

## TreeSHAP Algorithm

Efficient computation for tree ensembles:

- Exact Shapley values (not approximations!)
- Polynomial time:  $O(TLD^2)$  instead of exponential  $O(2^P)$
- Works for Random Forest, Gradient Boosting, XGBoost

# Why SHAP is Powerful for Finance

## Explainability is Crucial in Finance


- **Regulators:** Want to know why you made a trading decision
- **Clients:** Need to understand recommendations
- **Risk Management:** Must identify what drives predictions
- **Debugging:** Find if model learned something wrong

## Example Statement

*"The model recommended buying Bitcoin because momentum is strong (+1.2%), despite higher volatility (-0.8%)."*



# The Journey So Far



## Why Trees Are Perfect for Portfolios

- **Capture regime changes:** Bull vs. bear markets
- **Risk-return balance:** High return + low vol = great investment
- **Feature interactions:** Volatility × Momentum effects
- **Threshold effects:** Different behavior at vol > 30%

# Using Tree Predictions for Portfolio Weights

## Optimization Process

- 1 Train tree ensemble to predict asset returns and extract asset-specific predictions
- 2 Optimize weights using Sharpe ratios:

$$w_i = \frac{\exp(\text{Sharpe}_i)}{\sum_j \exp(\text{Sharpe}_j)}$$

- 3 Apply constraints:
  - $\sum_i w_i = 1$  (invest all money) and  $w_i \geq 0$  (no short selling)

## Example Allocation

- Bitcoin: +2% return, 3% volatility → 35% weight
- Ethereum: +1.5% return, 2% volatility → 40% weight
- Dogecoin: +3% return, 5% volatility → 25% weight

# What Are Hyperparameters?

## The Baking Analogy

Think of training a model like baking a cake:

- **Ingredients:** Your data (can't change)
- **Recipe:** Your algorithm (Random Forest, XGBoost)
- **Oven settings:** Your hyperparameters (temperature, time)

Just like 350°F for 30 min might be perfect, while 500°F for 10 min burns everything, **the right hyperparameters make a huge difference!**

# Grid Search vs Random Search

## Grid Search

**Strategy:** Test all combinations

- `n_estimators`: [50, 100, 200]
- `max_depth`: [5, 10, 15]
- `learning_rate`: [0.01, 0.1]

Total:  $3 \times 3 \times 2 = 18$  models

**Pro:** Guaranteed to find best

**Con:** Slow for many parameters

## Random Search

**Strategy:** Random sampling

- Try 50–100 random combinations
- More efficient!

**Surprising Result:** Often works as well as grid search but much faster!

**Why?** Usually only 2–3 hyperparameters really matter.

# Practical Tuning Tips

## Smart Tuning Strategy

### Round 1 – Coarse Search:

- Try wide range: depths [3, 10, 20], trees [50, 200, 500]
- Find general area: “Ah, depth around 10 seems best”

### Round 2 – Fine Search:

- Narrow down: depths [7, 10, 13], trees [150, 200, 250]
- Find exact sweet spot

## Key Rules

- 1 Always use cross-validation
- 2 Start with default values, then adjust
- 3 Don't spend hours tuning – diminishing returns!
- 4 More trees = almost always better (until you run out of time)

# Don't Trust Training Scores!

## Common Mistake

- Train model on January–October data
- Test on **same** January–October data
- Get 95% accuracy!
- **Problem:** Model memorized the answers!

This is like giving students the exam questions beforehand – of course they do well!

## Solution: Cross-Validation

- 1 Split data into 5 equal parts (folds)
- 2 Train on 4 folds, test on 1 fold
- 3 Repeat 5 times (each fold used as test once)
- 4 Average all 5 test scores → True performance!

# Comparing Random Forest vs XGBoost

## Example: 5-Fold Cross-Validation Results

	Random Forest	XGBoost
Fold 1	82%	87%
Fold 2	85%	88%
Fold 3	83%	86%
Fold 4	84%	87%
Fold 5	81%	87%
<b>Average</b>	83% $\pm$ 1.5%	87% $\pm$ 0.7%

**Winner:** XGBoost (higher score AND more consistent!)

**Statistical Testing:** Use t-test to verify difference is real ( $p\text{-value} < 0.05$ )

# What You Must Remember

- 1 **Decision trees** capture non-linear patterns but suffer from high variance
- 2 **Random Forest** (bagging) reduces variance through averaging decorrelated trees
- 3 **Gradient Boosting** learns sequentially from residual errors
- 4 **XGBoost** adds regularization and computational optimizations
- 5 **No feature scaling needed** – trees are scale-invariant
- 6 **Use SHAP** for explainability – don't treat trees as black boxes
- 7 **Hyperparameter tuning** is crucial for optimal performance
- 8 **Cross-validation** prevents overfitting and ensures generalization
- 9 Trees excel at **non-linear patterns** that linear models miss
- 10 **Ensembles**  $\gg$  single models for real-world finance



# Comparison Summary

Property	Random Forest	Gradient Boosting	XGBoost
Training	Parallel	Sequential	Sequential
Main Goal	↓ Variance	↓ Bias & Variance	Both + Speed
Tree Depth	Deep	Shallow (3–7)	Flexible
Learning Rate	N/A	Important	Important
Regularization	No	Limited	Yes (L1 + L2)
Speed	Fast	Moderate	Very Fast
Interpretability	Good	Good	Good
Competition Performance	Good	Better	<b>Best</b>

## Start Simple, Then Improve

- ① Train single decision tree – understand behavior
- ② Add Random Forest – should be much better!
- ③ Try Gradient Boosting and XGBoost – find best model
- ④ Analyze with SHAP – understand what model learned

## Warning Signs to Watch For

- Perfect training score (100%) → Overfitting!
- Big train/test gap (95% vs 70%) → Reduce complexity
- Weird feature importance → Check with SHAP
- Unstable predictions → Use more trees, cross-validation

# Quick Debugging Checklist

- ☐ Used cross-validation (not just train/test split)?
- ☐ Checked for overfitting (train vs test gap)?
- ☐ Tuned hyperparameters (tried different settings)?
- ☐ Analyzed feature importance (makes sense)?
- ☐ Created SHAP plots (understand decisions)?
- ☐ Compared multiple models (which is best)?
- ☐ Tested on completely unseen data (real performance)?

**If you can check all these boxes, you're doing great!**

## Example 1: Single Tree vs Random Forest

```
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import cross_val_score

# Single tree: High variance
tree = DecisionTreeRegressor(max_depth=10)
cv_scores = cross_val_score(tree, X, y, cv=5)
print(f"Tree CV R^2: {cv_scores.mean():.3f} +/- {cv_scores.std():.3f}")
# Output: 0.65 +/- 0.15 (high std!)

# Random Forest: Low variance
rf = RandomForestRegressor(n_estimators=100, max_depth=10)
cv_scores = cross_val_score(rf, X, y, cv=5)
print(f"RF CV R^2: {cv_scores.mean():.3f} +/- {cv_scores.std():.3f}")
# Output: 0.82 +/- 0.03 (low std!)
```

**Key Observation:** Random Forest dramatically reduces variance while maintaining accuracy!

## Example 2: XGBoost Regularization Impact

```
# No regularization: Overfits
xgb1 = XGBRegressor(reg_alpha=0, reg_lambda=0)
xgb1.fit(X_train, y_train)
print(f"Train R^2: {xgb1.score(X_train, y_train):.3f}") # 0.95
print(f"Test R^2: {xgb1.score(X_test, y_test):.3f}") # 0.70
# Gap of 0.25 indicates overfitting!

# With regularization: Generalizes better
xgb2 = XGBRegressor(reg_alpha=0.1, reg_lambda=1.5)
xgb2.fit(X_train, y_train)
print(f"Train R^2: {xgb2.score(X_train, y_train):.3f}") # 0.88
print(f"Test R^2: {xgb2.score(X_test, y_test):.3f}") # 0.85
# Gap of only 0.03 - much better generalization!
```

**Lesson:** Always use regularization with XGBoost!

# Integration with Previous Weeks

Week	Topic	Connection to Week 6
1	Data Collection	Raw price data → Features for trees
2	Data Cleaning	Handle missing values in tree inputs
3	Econometrics	Factor models = features for trees
4	Time Series	ARIMA/GARCH forecasts = tree inputs
5	ML Regularization	Ridge/Lasso = linear baseline; Trees = non-linear extension
6	<b>Tree Ensembles</b>	<b>Capture interactions &amp; non-linearities</b>
7	Dimensionality Reduction	PCA features → Trees (next week!)

## The Big Picture

- 1 Collect & clean data (Weeks 1-2)
- 2 Create features from time series (Weeks 3-4)
- 3 Try linear models first (Week 5)
- 4 Capture non-linearities with trees (Week 6)
- 5 Reduce dimensions if needed (Week 7)
- 6 Build optimal portfolio!

# Recommended Books

- ① **Hastie, T., Tibshirani, R., & Friedman, J. (2009).**  
*The Elements of Statistical Learning* (2nd ed.). Springer.  
Chapters 10, 15, 16.  
[The mathematical foundation – more advanced](#)
- ② **James, G., Witten, D., Hastie, T., & Tibshirani, R. (2023).**  
*An Introduction to Statistical Learning with Applications in Python* (2nd ed.). Springer.  
Chapter 8.  
[More accessible with Python examples](#)

## Free Resources

Both books are freely available online from the authors' websites!

# Key Research Papers

- ① **Breiman, L. (2001).** Random forests. *Machine Learning*, 45(1), 5–32.  
Original Random Forest paper
- ② **Friedman, J. H. (2001).** Greedy function approximation: A gradient boosting machine. *Annals of Statistics*, 29(5), 1189–1232.  
Original Gradient Boosting paper
- ③ **Chen, T., & Guestrin, C. (2016).** XGBoost: A scalable tree boosting system. *KDD 2016*, 785–794.  
XGBoost paper – most cited!
- ④ **Lundberg, S. M., & Lee, S.-I. (2017).** A unified approach to interpreting model predictions. *NeurIPS 2017*, 4765–4774.  
SHAP values paper



❶ **Gu, S., Kelly, B., & Xiu, D. (2020).**

Empirical asset pricing via machine learning.  
*Review of Financial Studies*, 33(5), 2223–2273.

Comprehensive comparison of ML methods for asset pricing, including tree ensembles. Shows XGBoost performs best for predicting stock returns.

❷ **Rasekhschaffe, K. C., & Jones, R. C. (2019).**

Machine learning for stock selection.  
*Financial Analysts Journal*, 75(3), 70–88.

Practical guide to using tree ensembles for portfolio construction and stock selection strategies.

# Statistical Testing for Model Comparison

## The Question

Is XGBoost **really** better than Random Forest, or did we just get lucky?

## Hypothesis Test

- $H_0$ : No difference between models
- $H_1$ : XGBoost performs better

Use **paired t-test** on cross-validation scores, if  $p\text{-value} < 0.05$ , the difference is statistically significant!

$$t = \frac{\bar{d}}{s_d / \sqrt{k}}$$

- $\bar{d}$  = mean difference in CV scores
- $s_d$  = standard deviation of differences
- $k$  = number of CV folds

# Statistical Testing in Python

```
from scipy import stats
from sklearn.model_selection import cross_val_score

# Get CV scores for both models
rf_scores = cross_val_score(rf_model, X, y, cv=5)
xgb_scores = cross_val_score(xgb_model, X, y, cv=5)

# Paired t-test
t_stat, p_value = stats.ttest_rel(xgb_scores, rf_scores)

print(f"RF scores: {rf_scores.mean():.3f} +/- {rf_scores.std():.3f}")
print(f"XGB scores: {xgb_scores.mean():.3f} +/- {xgb_scores.std():.3f}")
print(f"t-statistic: {t_stat:.3f}")
print(f"p-value: {p_value:.4f}")

if p_value < 0.05:
    print("Difference is statistically significant!")
else:
    print("No significant difference detected.")
```

# Next Week: Dimensionality Reduction

## Coming Next Week

**The Problem:** What if you have 100+ features? Trees can get confused!

**The Solution:**

- **PCA:** Compress 100 features into 10 “super-features”
- **Clustering:** Group similar assets together
- **Factor Models:** Find hidden patterns in markets

Thank you